

## Lab session #7 Artificial Neural Networks

Dr Zied M'NASRI

### Introduction

In this Lab session, we learn how to implement and train an Artificial Neural Networks to make predictions: a) binary classification, b) multiclass classification and c) regression. In this lab session, you will need to use:

1. Python tutorial (see Lab session 1 material)
2. Colab/Jupyter tutorial (see Lab session 2 material)
3. Lecture 7 notes
4. Lab 7 code example (Artificial Neural Networks with Backpropagation)
5. [Google's Colab](#) / Jupyter notebook

### Tutorial : Multi-class Classification using MLP with Error Backpropagation

Upload the *UB\_AI\_Lab7\_Example\_NN\_with\_BackPropagation.ipynb* to Google Colab and run it, to understand how ANN are trained with Backpropagation, using a random dataset.

In this tutorial, we use the functions provided in the code example

*UB\_AI\_Lecture7\_Example\_NN\_with\_BackPropagation.ipynb* to classify the example Iris dataset.

1. Upload the file *UB\_AI\_Lab7\_Example\_NN\_with\_BackPropagation.ipynb* To Google Colab or Jupyter Notebooks
2. Import the dataset Iris from sklearn.datasets

```
from sklearn import datasets
iris = datasets.load_iris()
print("Targets labels are: ",iris.target_names)
print("Target values are: ",np.unique(iris.target))
```

3. Separate Iris data into 2 data frames: X (data) and y (targets).

**Hint!** Use the method `pd.DataFrame()` and specify columns

```
X = pd.DataFrame(iris['data'],columns=['sepal length','sepal width','petal length','petal width'])
y = pd.DataFrame(iris['target'],columns=['target'])
print("Size of input features : ", X.shape)
print("Size of output targets : ", y.shape)
Nfeatures = X.shape[1]
print(f"Number of input features = ",Nfeatures)
Ntargets = y.shape[1]
print("Number of output targets = ",Ntargets)
```

4. Print the following results: Targets labels, Targets values, Size of input features, Size of output targets,

```
Targets labels are:  ['setosa' 'versicolor' 'virginica']
Target values are:  [0 1 2]
Size of input features :  (150, 4)
Size of output targets :  (150, 1)
Number of input features =  4
Number of output targets =  1
```

**Figure 1.1. Input and output data**

5. Split  $X$  and  $y$  into  $X\_train$ ,  $y\_train$ ,  $X\_test$  and  $y\_test$  using a split ratio 0.3 and `random_state= 42` using the method `train_test_split()`

**Hint!** You need to import `train_test_split()` from `sklearn.model_selection`

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=42)
for i in range(len(y_train)):
    y_train.iloc[i,0] = int(y_train.iloc[i,0])
```

6. Form a list called `dataset_train` by concatenating  $X\_train$  and  $y\_train$ , and print the first 5 rows of `dataset_train`

**Hint!** Use the method `pd.concat()` with its attribute method `values.tolist()`

```
dataset_train = pd.concat([X_train, y_train],
axis=1).values.tolist()
print(f"dataset_train: {dataset_train[0:5]}")
```

7. Do the same for dataset\_test with X\_test and Y\_test

```
dataset_test = pd.concat([X_test, y_test], axis=1).values.tolist()
print(f"dataset_test: {dataset_test[0:5]}")
```

```
dataset_train: [[4.6, 3.6, 1.0, 0.2, 0.0], [5.7, 4.4, 1.5, 0.4, 0.0], [6.7, 3.1, 4.4, 1.4, 1.0], [4.8, 3.4, 1.6, 0.2, 0.0], [4.4, 3.2, 1.3, 0.2, 0.0]]
dataset test: [[6.1, 2.8, 4.7, 1.2, 1.0], [5.7, 3.8, 1.7, 0.3, 0.0], [7.7, 2.6, 6.9, 2.3, 2.0], [6.0, 2.9, 4.5, 1.5, 1.0], [6.8, 2.8, 4.8, 1.4, 1.0]]
```

**Figure 1.2. First 5 rows of the training and test sets**

8. In Section “Network initialisation”:

- Input the number of nodes in the hidden layer
- Set n\_inputs as the number of features as the length of X\_train[0]
- Set n\_outputs as the length of the value counts in y\_train
- Hint!** Use the attribute method .value\_count()
- Initialize the network by calling the function initialize\_network() (defined in the example)
- Print the initialized network

```
# Network initialisation
n_hidden = int(input('Enter the number of nodes in the hidden
layer: '))
n_inputs = len(dataset_train[0]) - 1
n_outputs = len(y_train.value_counts())
print(f"number of input features = {n_inputs}")
print(f"number of output labels = {n_outputs}")
print(f"number of nodes in the hidden layer = {n_hidden}")
network = initialize_network(n_inputs, n_hidden, n_outputs)
print(f"Initial network weights: {network}")
```

```
Enter the number of nodes in the hidden layer: 2
number of input features = 4
number of output labels = 3
number of nodes in the hidden layer = 2
[{'weights': [0.4896935204622582, 0.029574963966907064, 0.04348729035652743, 0.703382088603836, 0.9831877173096739]}]
```

**Figure 1.3. Network initialization**

9. In Section “Network training”:

- Enter the learning rate (between 0 and 1)
- Enter the number of epochs
- Print the parameters above, plus the initial weights
- Train the networks using the function `train_network()` mentioned in the example code

```
# Network training
sum_error = 0
l_rate = float(input('Enter the learning rate (between 0 and 1): '))
n_epoch = int(input('Enter the number of epochs: '))
print(f"Selected learning rate = {l_rate}")
print(f"Selected number of epochs = {n_epoch}")
train_network(network, dataset_train, l_rate, n_epoch, n_outputs)
```

10. In section “Test making predictions with the network”:

- Replace the test dataset by `dataset_test` (created above)
- Print the target and the predicted values

```
# Test making predictions with the network
print(f"final networks: {network}")
target = prediction = np.zeros(len(dataset_test))
for i in range(len(dataset_test)):
    row = dataset_test[i][:]
    print(row)
    target[i]=int(dataset_test[i][-1])
    prediction[i] = predict(network, row)

print("targets= ", target)
print("predictions= ", prediction)
targets= [2. 0. 2. 2. 0. 1. 2. 2. 2. 0. 0. 0. 0. 2. 2. 2. 2. 2. 0. 2. 0. 2.
2. 2. 2. 0. 0.]
predictions= [2. 0. 2. 2. 0. 1. 2. 2. 2. 0. 0. 0. 0. 2. 2. 2. 2. 2. 0. 2. 0. 2.
2. 2. 2. 0. 0.]
```

**Figure 1.4. Targets and predictions after test**

11. In the section “Evaluate NN model”, use the method `classification_report()` to calculate the evaluation metrics. Print the classification report

```
# Evaluate NN model
#from sklearn import metrics
from sklearn.metrics import confusion_matrix, classification_report
print(f"Confusion matrix: \n {confusion_matrix(target,
prediction)}")
print(f"Classification_report: \n {classification_report(target,
prediction)}")
```

Confusion matrix:				
[[10  0  0]				
[ 0  1  0]				
[ 0  0 19]]				
Classification_report:				
	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	10
1.0	1.00	1.00	1.00	1
2.0	1.00	1.00	1.00	19
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

**Figure 1.5. Evaluation metrics on the test set**

## Exercises

### Exercise #1- Binary classification using MLP with predefined functions and grid search

In this exercise, we aim at using predefined functions in Python libraries (Scikit learn and Pandas) to: a) Preprocess data, b) Initialize and train an MLP, c) Improve the evaluation metrics using a grid search to select the best performing parameters.

1. Upload the file *play\_tennis.csv* (see lab 5 or lab 6) and print the dataset header

	Day	Outlook	Temperature	Humidity	Wind	PlayTennis
0	D1	Sunny	85	85	Weak	No
1	D2	Sunny	80	90	Strong	Yes
2	D3	Overcast	83	78	Weak	Yes
3	D4	Rain	70	96	Weak	Yes
4	D5	Rain	68	80	Weak	No

**Figure 2.1. play\_tennis dataset header**

2. Replace the values of the features Outlook and Wind to numerical, as follows (see lab 6):  
 Outlook → Sunny = 0, Overcast = 0.5, Rain = 1  
 Wind → Weak = 0, Strong = 1  
 PlayTennis → No = 0, Yes = 1
3. Normalise the values of Temperature using the min-max normalization  
**Hint!** Use the following formula:

```
dataset['featureName'] = (dataset['featureName'] - min(dataset['featureName'])) / (max(dataset['featureName']) - min(dataset['featureName']))
```

	Day	Outlook	Temperature	Humidity	Wind	PlayTennis
0	D1	0.0	0.833333	0.738095	0	0
1	D2	0.0	0.714286	0.857143	1	1
2	D3	0.5	0.785714	0.571429	0	1
3	D4	1.0	0.476190	1.000000	0	1
4	D5	1.0	0.428571	0.619048	0	0

**Figure 2.2. Normalized dataset**

4. Declare X and y as the feature columns and the target column, respectively

**Hint!** Use the attribute `datasetName.iloc` (see Lab 6)

- Split X and y into X\_train, X\_test, y\_train, y\_test using a test/train ratio = 0.3 (and <0.5 in any case). Use `random_state = 42`. Print the shape of each subset.

```
size of training features : (279, 4)
size of training targets : (279,)
size of test features : (120, 4)
size of test targets : (120,)
```

**Figure 2.3. Training and test subsets**

- Define an MLP classifier using the method `MLPClassifier()` with the following parameters:

```
hidden_layer_sizes=(200, 100),
max_iter = 300,
activation = 'relu',
solver = 'sgd',
learning_rate='constant',
learning_rate_init=0.001,
batch_size='auto',
tol=0.00001,
verbose=True
```

**Hint!** Import `MLPClassifier` from `sklearn.neural_network`

- Train the MLP classifier and print the trained weights

**Hint!** Use the method `MLPClassifierName.fit()` to train the classifier.

The weights are stored in `MLPClassifierName.coefs_`

```
MLP classifier weights = [array([[ 0.07255927, -0.02182532, -0.05824862,  0.07662956,  0.09351607,
   -0.12959692, -0.12025699, -0.16760501,  0.06665141,  0.09019231,
   0.12430765, -0.00388537,  0.12044874,  0.12309126, -0.06039614,
   -0.08944968,  0.11498595, -0.0411311 , -0.054184 , -0.13232408,
   -0.13813591,  0.09055127, -0.04877255,  0.11799414,  0.10812991,
   0.06079494,  0.05857149,  0.0789826 ,  0.15238499,  0.11716647,
```

**Figure 2.4. MLP classifier weights**

- Test the trained model `MLPClassifierName` on `X_test`

**Hint!** Use the method `MLPClassifierName.predict()`

- Calculate the accuracy and print the classification report


**Hint!** Import and use the methods `accuracy_score()` and `classification_report()` from `sklearn.metrics`

	precision	recall	f1-score	support
0	0.00	0.00	0.00	55
1	0.54	1.00	0.70	65
accuracy			0.54	120
macro avg	0.27	0.50	0.35	120
weighted avg	0.29	0.54	0.38	120

**Figure 2.5. Classification report**

10. Plot the loss curve

**Hint!** The loss curve is stored in `MLPClassifierName.loss_curve_`



**Figure 2.6. Loss curve**

11. To improve the performance of the MLP classifier, we can try different configurations using a grid search. Define the grid parameters as follows:

```
param_grid = {'hidden_layer_sizes': [(150,100,50), (120,80,40),
(100,50,30)], 'max_iter': [50, 100, 150], 'activation': ['tanh', 'relu'],
'solver': ['sgd', 'adam'], 'alpha': [0.0001, 0.05], 'learning_rate': ['constant', 'adaptive'], }
```

12. Define a `gridName` model using the method `GridSearchCV()` with the following parameters:

`N_jobs = -1, cv = 5`

13. Train the `gridName` model with `X_train` and `y_train`

14. Print the `gridName` model best parameters, found after training

**Hint!** The best parameters are stored in `gridName.best_parameters_`

15. Make predictions on the test set using the `gridName` model.

16. Print the accuracy and the classification report. What do you notice about the models performance?

Accuracy: 0.78					
	precision	recall	f1-score	support	
0	0.77	0.73	0.75	55	
1	0.78	0.82	0.80	65	
accuracy			0.78	120	
macro avg	0.77	0.77	0.77	120	
weighted avg	0.77	0.78	0.77	120	

**Figure 2.7. Classification report after parameters optimisation using grid search**

## Exercise #2- Regression with MLP and input data standardization

This exercise aims at using MLP for regression, i.e. continuous-valued output prediction. We use also data standardization, i.e. transforming the input data so that it has a zero-mean and a unit-variance, instead of min-max normalisation. Also we use a different evaluation metric, i.e. the root mean square error, since the output is not categorical.

1. Import the following packages:

```
import numpy as np
from sklearn import datasets
from sklearn.datasets import fetch_california_housing
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_squared_error
```

2. Declare housing as the result of `fetch_california_housing()` and X, y as the attributes `housing.data()` and `housing.target()`, respectively. Print X and y and their size.

```
Size of X: (20640, 8)
[[ 8.3252    41.      6.98412698 ...  2.55555556
  37.88   -122.23     ]  6.23813708 ...  2.10984183
 [ 8.3014    21.      6.28813559 ...  2.80225989
  37.86   -122.22     ] ...
 [ 7.2574    52.      8.28813559 ...  2.80225989
  37.85   -122.24     ]
...
[ 1.7       17.      5.20554273 ...  2.3256351
 39.43   -121.22     ]  5.32951289 ...  2.12320917
 [ 1.8672    18.      5.25471698 ...  2.61698113
 39.43   -121.32     ] ...
[ 2.3886    16.      5.25471698 ...  2.61698113
 39.37   -121.24     ]]
Size of y: (20640,)
[4.526 3.585 3.521 ... 0.923 0.847 0.894]
```

**Figure 3.1. Data and targets and the size of each**

3. Transform X and y to X\_std and y\_std using the mean-variance standardisation formula:

$$z_{std} = \frac{z - \mu}{\sigma}$$

Where  $\mu$  is the mean value of (z) and  $\sigma$  is its standard deviation

**Hint!** For each of X and y, mean and standard deviation are obtained using the attribute methods `np.mean(variableName, axis = 0)` and `np.std(variableName, axis = 0)`

4. To visualize the effect of standardisation, plot the histograms of  $y$  and  $y_{std}$

Hint! Use the following methods:


```
sns.set(color_codes=True)
sns.set_style("whitegrid")
```

then for each variable:

```
sns.histplot(variableName)
```

Finally

```
plot.show()
```



**Figure 3.2. Histograms of original (blue) and standardised (amber) targets**

5. Split each of  $X_{std}$  and  $y_{std}$  into training and test sets using a test/training ratio = 0.3 and `random_state = 1`. Print the size of each subset

```
Size of X_std_train: (14448, 8)
Size of X_std_test: (6192, 8)
Size of y_std_train: (14448,)
Size of y_std_test: (6192,)
```

**Figure 3.3. Size of each split**

6. Initialize an MLP regressor using the following parameters:

```
activation='relu',
hidden_layer_sizes=(10, 100),
alpha=0.001,
random_state=20,
early_stopping=False
```


**Hint!** Use the built-in function `MLPRegressor()`

7. Train the model using `X_std_train` and `y_std_train`

**Hint!** Use the built-in function `mlpRegressorName.fit()`

8. Plot the training loss curve

**Hint!** When calling the function `plt.plot()`, use the attribute `mlpRegressorName.loss_curve_`



**Figure 3.4. Loss curve on the training set**

9. Make predictions on the test set using `X_std_test`

**Hint!** Use the function `mlpRegressorName.predict()`

10. Transform the predicted values into the real range of the targets

**Hint!** Use the reverse transform, i.e `z = z_std*np.std(z)+np.mean(z)`

```
Standardised predictions: [ 1.56297565 -1.24122407  0.28785319 ... -0.76120006 -0.41375626
```

```
 0.14063342]
```

```
Transformed predictions: [3.87211986 0.63627471 2.40072008 ... 1.19018795 1.59111315 2.23083904]
```

**Figure 3.5. Standardised and transformed predictions**

11. To evaluate the predicted values, calculate the Root Mean Square Error (RMSE), and its relative value (Relative RMSE), between the actual targets (`y_test`) and the transformed predictions.

**Hint!** Use the function `mean_squared_error()` to calculate MSE and `np.sqrt()` for the square root.